

# 1 Evaluation

## 1.1 Experimental Design

Our evaluation is based on two benchmarks with significantly different access patterns. The first is quicksort (Qsort). This benchmark first allocates a large array of random numbers, and then sorts it using the well-known quicksort algorithm. Quicksort is a divide and conquer algorithm that automatically partitions the input array into small local blocks before performing a final sort.

**Code-friendly description of raw data involved in this experiment (defines to the user what types of questions are possible). This field is collapsable. It is unclear if this should be a real piece of code or an abstract description.**

The other benchmark is a de novo genome assembly benchmark (Gen). Gen begins by loading a large text file that represents raw genome data. Raw genome data consists of short, overlapping, sequences of base-pairs called "contigs", the goal is to align these overlapping contigs into a single contiguous sequence representing a genome. This is done by loading contigs into a large hash table and probing into it repeatedly to find matching sequences. This leads to very little locality and unpredictable access patterns. Furthermore, Gen performs file I/O on the input which allows for more complex OS interactions.

### ▼ Schema

```
# Data Schema
result = NamedTuple(
    'mean', # Mean results from 10 runs
    'std', # Standard deviation from 10 runs
)

# mean and std have the same set of keys
result.mean.keys
[ 't_run', 't_bookkeeping', 't_rmem_write',
  't_rmem_read', 'n_fault', 't_fault',
  'n_swapfault', 'n_pfa_fault', 'n_early_newq',
  'n_evicted', 'n_fetched', 'n_kpfad',
  't_kpfad', 'slowdown' ]

# Datasets
# Paging to remote memory in SW ('baseline')
base_res = ingest_run('raw/results_baseline.csv', 'rv')

# Paging using the PFA to real memblade
pfa_res = ingest_run('raw/results_pfa.csv', 'rv')
```

## 2 End-to-End Performance

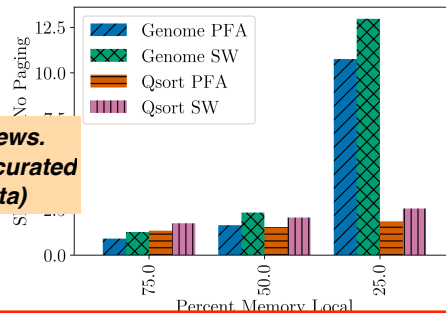
Both benchmarks were both run under a cgroup in Linux in order to reduce the need to share administrators use today to control application memory consumption (e.g., in containers). In this experiment, we disable kpfad in order to isolate the batching of new-page management from the scheduling flexibility offered by kswapd's asynchrony. The PFA was configured to allow up to 64 outstanding page faults before book-keeping was performed.

### ▼ Linux Config

```
CONFIG_PFA=y
CONFIG_PFA_EM=n
CONFIG_MEMBLADE_EM=n
CONFIG_PFA_VERBOSE=n
CONFIG_PFA_DEBUG=n
CONFIG_PFA_PFLAT=n
CONFIG_PFA_KPFAD=n
CONFIG_PFA_FREEQ_SIZE=64
CONFIG_PFA_NEWQ_SIZE=64
CONFIG_PFA_EVICTQ_SIZE=1
```

Live updated data from experimental results  
(uses tangle library from ExplainableExplanations)

Both applications use 64MB of memory at their peak. We then varied the cgroup memory limit from 100% (64MB) down to 25% (16MB), triggering increasing levels of paging. For both benchmarks, the PFA reduces end to end run time by up to 1.4x.



Tabbed interface to alternative data views.  
"Original" view is static and immutable, curated  
by author (still derived from base data)

Original User Explore

Figure 1: PFA vs Baseline without kswapd. Applications run approximately 20-40% faster when the PFA is enabled.

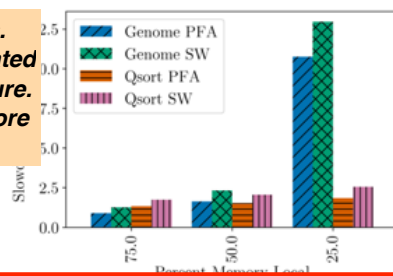
5

**Tabbed interface to alternative data views.**  
**“User” interface shows the code that generated the figure and allows users to tweak the figure.**  
**This constrained question space may be more approachable to less experienced users.**

```
speedup = base_res.mean['t_run'] /  
pfa_res.mean['t_run']
```

```
ax = speedup.plot.bar()  
ax.set_ylabel('% Improvement')  
ax.set_xlabel('Percent Memory Local')
```

```
plt.show()
```



**Original** **User** **Explore**

Figure 2: PFA vs Baseline without kswapd. Applications run approximately 20-40% faster when the PFA is enabled.

6

**Tabbed interface to alternative data views.**  
**“Explore” is a raw, programmatic view of the data that allows readers to ask open-ended questions**

```
In [7]: pfa_res['python'].mean
```

	okkeeping	t_rmem_write	t_rmem_read	n_fault	t_fault	n_swapfault
	0.005289	0.002702	0.0	10656.333333	0.079287	0.0
75.0	1.730389	0.031244	0.009852	0.0	17876.000000	0.105684
50.0	1.952219	0.137448	0.032581	0.0	36312.666667	0.259532
25.0	3.997994	0.843604	0.179366	0.0	190521.333333	1.905531

**Original** **User** **Explore**

Figure 3: PFA vs Baseline without kswapd. Applications run approximately 20-40% faster when the PFA is enabled.